

# High Performance Computing for Genomics

Part I: High Performance Computing

# Overview

- Login
- Storage
- Credit system
- Module system
- Hardware
- PBS
- Job Submission

# Login

Login is possible through various methods.

Command line example (Linux/Mac):

```
$ ssh -X vsc3XXXX@login1-tier2.hpc.kuleuven.be|
```

Or via putty (Windows)

# List of Login nodes

Node	URL
Genius	login1-tier2.hpc.kuleuven.be
Genius	login2-tier2.hpc.kuleuven.be
Genius (NX Client)	login3-tier2.hpc.kuleuven.be
Genius (NX Client)	login4-tier2.hpc.kuleuven.be

ssh -X [vsc3XXXX@login1-tier2.hpc.kuleuven.be](https://login1-tier2.hpc.kuleuven.be)

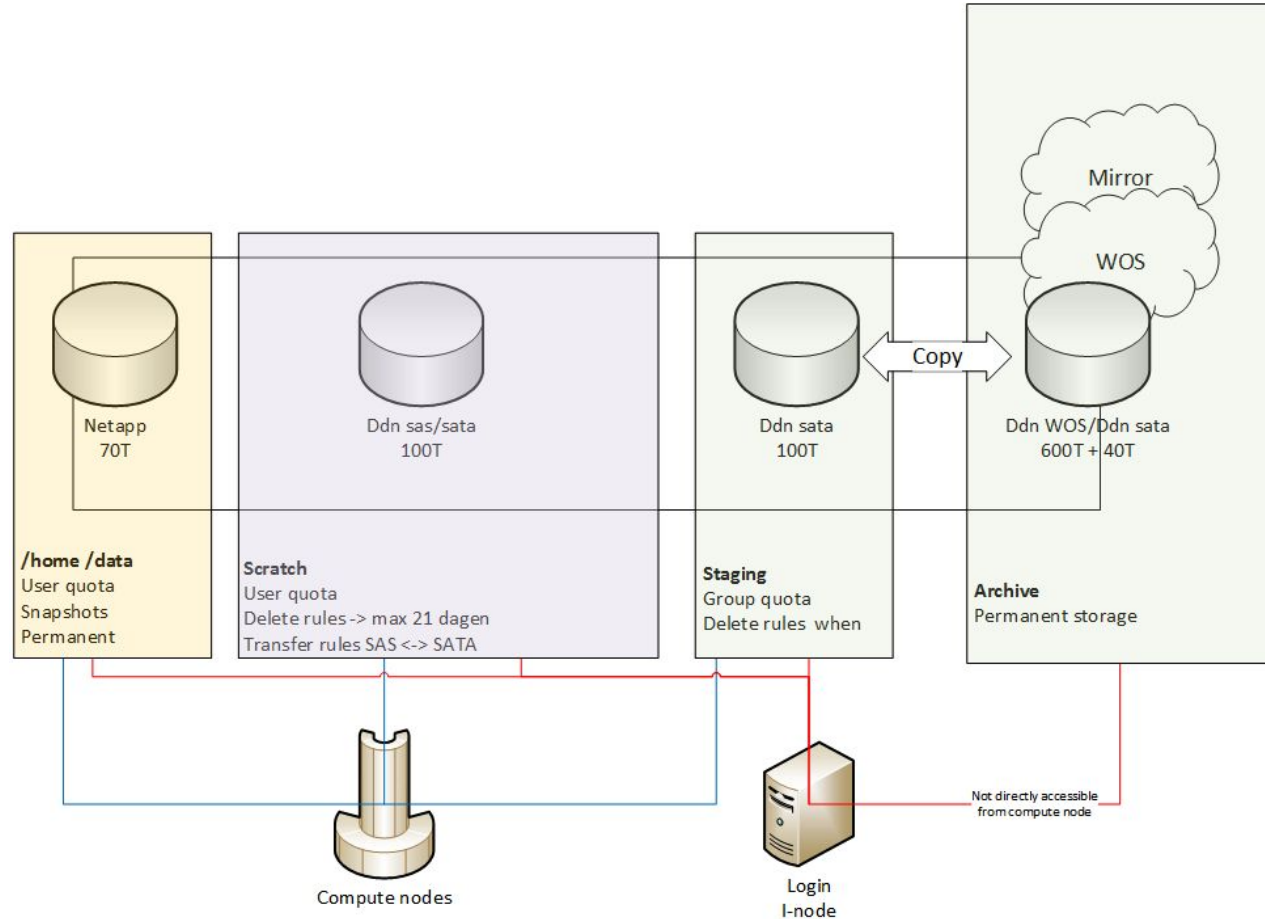
If you are not logged in yet,  
Now is the moment

# Overview

- Login
- Storage
- Credit system
- Module system
- Hardware
- PBS
- Job Submission

# Storage

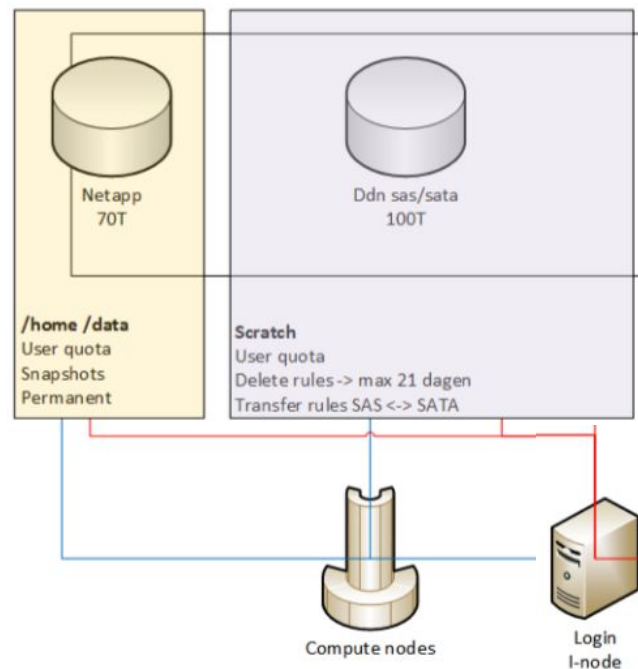
- NAS storage, fully back-up with snapshots for /home and /data
- Scratch storage, fast parallel filesystem
- Staging storage, to store current working projects
- Archive storage, to store large amounts of data for long time



# Personal Storage

Personal Storage, only accessible by the owner  
Fully Backup

- Home Directory
  - \$VSC\_HOME or /user/leuven/3XX/vsc3XXXX
  - 25GB
  - Important data, configuration files
- Data Directory
  - \$VSC\_DATA or /data/leuven/3XX/vsc3XXXX
  - 75GB
  - Important data, bigger then Home
- Scratch Space
  - \$VSC\_SCRATCH or /scratch/leuven/3XX/vsc3XXXX
  - 100GB
  - Temporary data, will be deleted within 21 days

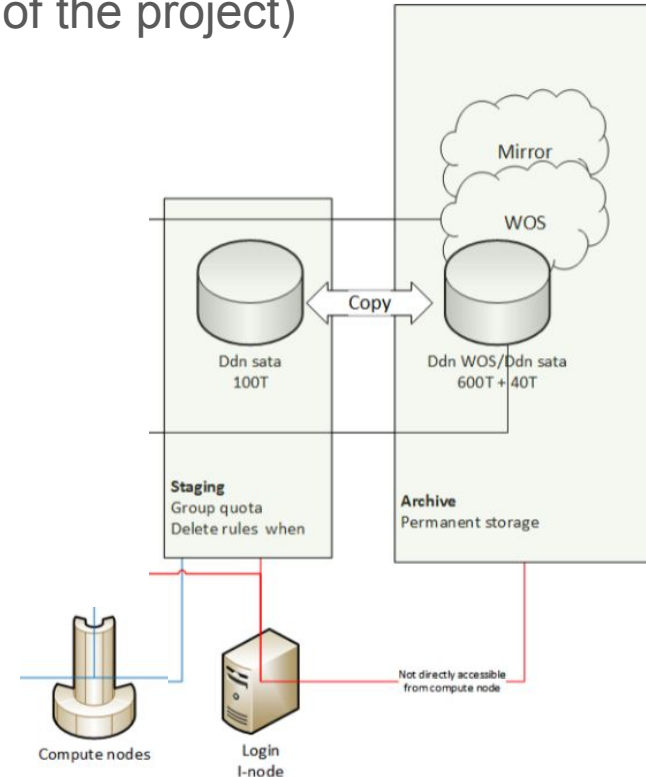




# Project Storage

Permissions are managed per group (usually the group of the project)


- **Staging Space**
  - Storage for project files (while working at the project)
  - Size is project dependable (per 1TB)
- **Archive Storage**
  - Long term storage and backup for project files (or staging)
  - Size is project dependable (per 1TB)
  - Only accessible from the login node



# Temporary Storage

Scratch: fast storage, used during the processing of the data

- Node Scratch
  - `$VSC_SCRATCH_NODE`
  - Size is machine dependable, min 150 GB
  - Is NOT accessible from the login node
- Site Scratch
  - `$VSC_SCRATCH_SITE`
  - Similar as `$VSC_SCRATCH_NODE`
- Personal Scratch
  - `$VSC_SCRATCH`
  - Personal scratch, 100GB
  - Accessible from the login node
  - Data will be deleted within 21 days



# Exercise 1

Find the path for the following locations:

- \$VSC\_HOME
- \$VSC\_DATA
- \$VSC\_SCRATCH
- \$VSC\_SCRATCH\_NODE

# Copy Data

```
$ rsync -ahr --progress /staging/leuven/stg_00019/workshop/* .
```

`rsync --progress -ahr filename targetDirectory`

<code>--progress</code>	Show the progress of the copy
<code>-a</code>	Preserve almost everything
<code>-h</code>	Output numbers must be Human readable
<code>-r</code>	Copy directories recursively

# Exercise 1b

Go to your data directory

Clone the github repository:

```
$ git clone https://github.com/GenomicsCoreLeuven/vsc_ngs_workshop.git
```

# Overview

- Login
- Storage
- Credit system
- Module system
- Hardware
- PBS
- Job Submission

# Credit System

VSC ~ regular bank

Each job costs a certain number of credits

Each job belongs to a project

Each project account has its own credits

Users have access to one or multiple project accounts  
(multiple users can have access to the same project account)



# Credit System: balance

Load the accounting module

```
$ module load accounting
```

Check the amount of credits at your disposal

```
$ mam-balance
```



# Credit System: statement

Get an overview of the transactions

```
$ mam-statement
```

Get an overview of the transactions of a certain project, over a certain time

```
$ mam-statement -g lp_projectname -s 2015-09-01 -e 2015-09-30
```

-g	groupname
-u	username
-s	start
-e	end

# Credits System: quotes

Estimating the cost of a job is possible by requesting a quote:

```
$ gquote -l nodes=3:ppn=4:ivybridge,pmem=2gb,walltime=48:00:00
```

- nodes=3:ppn=4:ivybridge  
3 nodes, with 4 processors per node of the ivybridge type
- pmem=2gb  
2GB of memory (RAM)
- walltime=48:00:00  
The job will run for maximum 48 hours

# Job Cost Calculation

The effective cost of the job is calculated using this formula:

$$(0.000278 * \text{nodes} * \text{walltime} + \text{startup}) * \text{nodetype}$$

nodes	The number of nodes that were reserved
walltime	The effective duration of the job (in seconds)
startup	0.1, it is added for the overhead of scheduling
nodetype	A representation of the node type's performance

## Exercise 2

- Get a quote for 5 minutes computing on 1 node, 5 processors of the ivybridge type
- Get a quote for 20 minutes computing on 2 nodes, 4 processors of the haswell type, using 100Gb of memory
- Get a quote for 1 hour for both the ivybridge and the haswell processor

# Overview

- Login
- Storage
- Credit system
- Module system
- Hardware
- PBS
- Job Submission

# Module System

For many **programs**, **multiple versions** are installed, and each version requires **specific libraries**.

**Toolchains** consists of **compilers and libraries**, together with **software** depending on those libraries.

The **module** system inside the toolchain is used to manage the **environment variables**, and all **dependencies** to resolve possible conflicts.

# Module System: Module av

Many software packages are installed as modules:

```
$ module av
```

```
----- /apps/leuven/skylake/2018a/modules/all -----
```

```
7zip/16.02-GCCcore-6.4.0
```

```
ABINIT/8.4.4-intel-2018a
```

```
ABINIT/8.8.3-intel-2018a
```

```
(D)
```

```
ABYSS/2.0.2-foss-2018a
```

```
ADMIXTOOLS/5.1-foss-2018a
```

```
ANTLR/2.7.7-intel-2018a-Python-3.6.4
```

```
ANTs/2.3.1-foss-2018a-Python-2.7.14
```

```
ATK/2.28.1-foss-2018a
```

# Module System: load and unload

Load a module

```
$ module load BEAST
```

```
$ module load BEAST/2.1.2
```

```
$ module load BEAST/2.1.2 zlib/1.2.8-foss-2014a
```

Unload a module

```
$ module unload BEAST
```

Unload all modules

```
$ module purge
```



# Module System: list

## View all loaded modules

```
$ module list
```

```
Currently Loaded Modulefiles:
```

- 1) /thinking/2014a
- 2) Java/1.7.0\_51
- 3) icc/2013.5.192
- 4) ifort/2013.5.192
- 5) impi/4.1.3.045
- 6) imkl/11.1.1.106
- 7) intel/2014a
- 8) beagle-lib/20140304-intel-2014a
- 9) BEAST/2.1.2
- 10) GCC/4.8.2
- 11) OpenMPI/1.6.5-GCC-4.8.2
- 12) gomp/2014a
- 13) OpenBLAS/0.2.8-gomp-2014a-LAPACK-3.5.0
- 14) FFTW/3.3.3-gomp-2014a
- 15) ScaLAPACK/2.0.2-gomp-2014a-OpenBLAS-0.2.8-LAPACK-3.5.0
- 16) foss/2014a
- 17) zlib/1.2.8-foss-2014a

# Overview

- Login
- Storage
- Credit system
- Module system
- Hardware
- PBS
- Job Submission

# Hardware

- Genius  
Thin node cluster
- Superdome  
Shared Memory Processing
- Accelerators  
Like GPUs
- Storage

# Hardware: Genius

- Genius
  - Main Cluster
  - 3 types of nodes:
    - Thin nodes (Skylake)  
86 nodes, 2x18 cores/node, 192GB RAM
    - Big nodes (Skylake)  
10 nodes, 2x18 cores/node, 768 GB RAM  
Partition: bigmem
    - GPGPU (Skylake)  
20 nodes, 2x18 cores/node, 192GB RAM  
4 NVIDIA P100, 16GB GDDR  
Partition: gpu
- Superdome

# Hardware: Genius

- Genius
- Superdome
  - 1 type of node:
    - Shared Memory nodes (Skylake)  
8 nodes, 14 cores/node, 750GB RAM  
Partition: superdome
  - Up to 6TB RAM!

# Hardware: Overview

Cluster	Partition	#cores (threads) per node	Usable Memory	#Credits/hour
Genius		36	192GB	10
Genius	bigmem	36	768GB	10
Genius	gpu	36 + 4 GPU	192GB	
Genius	superdome	14	Shared 750GB	10

The Operating System also needs some RAM, good practice is to reserve 4GB.

# Hardware: which task on which cluster?

Cluster	Partition	Task Description	Task Example
Genius	Regular And gpu	Memory low jobs, with lots of I/O	<ul style="list-style-type: none"><li>• Alignment</li><li>• Read Mapping</li><li>• Variant Calling</li><li>• Read Counting</li><li>• ...</li></ul>
Genius	bigmem	High memory jobs	<ul style="list-style-type: none"><li>• Full Genome (elprep, ...)</li><li>• Small De Novo Assemblies</li><li>• Reference-based Assemblies</li><li>• ...</li></ul>
Genius	superdome	Really High memory jobs, computing power less important	<ul style="list-style-type: none"><li>• Large De Novo Assemblies</li><li>• ...</li></ul>

# Overview

- Login
- Storage
- Credit system
- Module system
- Hardware
- PBS
- Job Submission



# Portable Batch System (PBS)

Portable Batch System (or simply PBS ) is the name of computer [software](#) that performs job scheduling. Its primary task is to [allocate computational tasks](#), i.e., batch jobs, among the available computing resources. It is often used in conjunction with UNIX [cluster environments](#). PBS is supported as a [job scheduler](#) mechanism by several meta schedulers ...

-- by Wikipedia

# PBS script

Every shell script (.sh) can be turned into a PBS script (.pbs).

A typical scripts will need the following adjustments:

1. Include PBS headers
2. Load the software modules
3. Optimize I/O using a scratch node
4. Check if results are copied/stored in a data or staging directory

# PBS header

Headers contain the run parameters

Each line starts with #PBS then the parameter and the value

```
#!/bin/bash -l
#PBS -l walltime=12:00:00
#PBS -l mem=100gb
#PBS -l nodes=1:ppn=20
#PBS -M mail@mail.com
#PBS -m aeb
#PBS -N jobname
#PBS -A lp_projectname
```

# PBS header: job description

Walltime: the maximum time the job can run

```
#PBS -l walltime=12:00:00
```

Mem: the maximum memory available for the job

```
#PBS -l mem=100gb
```

Number of nodes, processors per node (ppn) and the processor type

```
#PBS -l nodes=1:ppn=20
```

# PBS header: notifications

You can send a mail @ certain conditions

-M specifies the mail adres

```
#PBS -M mail@mail.com
```

```
#PBS -m aeb
```

m the conditions when a mail has to be send:

b	When the job begins
e	When the job ends
a	When the job is aborted

# PBS header: billing

The name of the job, this will be visible for other users

```
#PBS -N jobname
```

The project that will be used for the billing

```
#PBS -A lp_projectname
```

# PBS script

Every shell script (.sh) can be turned into a PBS script (.pbs).

A typical scripts will need the following adjustments:

1. Include PBS headers
2. **Load the software modules**

```
module load BEAST/2.1.2
```

3. Optimize I/O using a scratch node
4. Check if results are copied/stored in a data or staging directory

# PBS script

Every shell script (.sh) can be turned into a PBS script (.pbs).

A typical scripts will need the following adjustments:

1. Include PBS headers
2. Load the software modules
3. **Optimize I/O using a scratch node**  
**`$VSC_SCRATCH`, `$VSC_SCRATCH_NODE`**  
**Scratch is fast temporary storage**
4. Check if results are copied/stored in a data or staging directory



# PBS script

Every shell script (.sh) can be turned into a PBS script (.pbs).

A typical scripts will need the following adjustments:

1. Include PBS headers
2. Load the software modules
3. Optimize I/O using a scratch node
4. **Check if results are copied/stored in a data or staging directory**  
**Scratch is temporary, when the job is finished, data will be lost**

# Overview

- Login
- Storage
- Credit system
- Module system
- Hardware
- PBS
- Job Submission

# Job submission

On all partitions:

```
$ qsub run-job.pbs
```

You have to specify the partition in the PBS script (superdome request an extra line):

```
#PBS -L tasks=1:lprocs=14:place=numanode:memory=700gb  
#PBS -l partition=superdome
```

When a job is submitted, the job id is returned.  
jobIDs are unique.

# Jobs progress

An overview of all your jobs:

```
$ qstat -u vsc3XXXX  
hpc-p-svcs-5:
```

Req'd		Req'd		Elap		Queue	Jobname
Job ID				Username			
SessID	NDS	TSK	Memory	Time	S		Time
-----							
-----							
30017323.hpc-p-svcs-5				vsc3XXXX		q21d	jobname
197071	1	10	1gb	480:00:00	R		00:04:16

# Job estimated start

Get the estimated start of a job:

```
$ showstart 30017323.hpc-p-svcs-5
job 30017323 requires 10 procs for 20:00:00:00
Estimated Rsv based start in 00:00:00 on Mon Apr
18 16:29:10
Estimated Rsv based completion in 20:00:00:00 on Sun May
8 16:29:10
Best Partition: smp1
NOTE: job is running
```

# Job overview

An overview of every parameter and all used resources can be requested:

```
$ checkjob 30017323.hpc-p-svcs-5
job 30017323
AName: jobname
State: Running
Creds:  user:vsc3XXXX  group:vsc3XXXX  account:lp_projectname
class:q21d  qos:normal
WallTime:   00:02:43 of 20:00:00:00
SubmitTime: Mon Apr 18 16:28:56
  (Time Queued  Total: 00:00:14  Eligible: 00:00:14)
StartTime: Mon Apr 18 16:29:10
TemplateSets:  DEFAULT
Total Requested Tasks: 10
Req[0]  TaskCount: 10  Partition: smp1
Memory >= 1024M  Disk >= 0  Swap >= 0
Dedicated Resources Per Task: PROCS: 1  MEM: 24G
```

# Job stop/kill/delete

Sometimes a job has to be killed.

```
qdel 30017323.hpc-p-svcs-5
```

# Exercise 3

1. Open paths.pbs  
Edit the PBS header (change the mail, and accounting information)
2. Ask a quote based on the PBS header
3. Start the job on thinking
4. Check the start time and status of the job
5. Check the output of the job



# Coffee break

